

Advanced Concepts

Special Pod Cases. Autoscaling and Scheduling

Daemon Sets and Jobs. Ingress




kubernetes

SoftUni Team

Technical Trainers

 Software University



SoftUni



Software University

<https://softuni.bg/>

Have a Question?




sli.do

#Kubernetes

[facebook.com
/groups/kubernetesnovember2025](https://facebook.com/groups/kubernetesnovember2025)

Homework Progress





Previous Module (M4)

Quick overview

Table of Contents

1. (Persistent) Volumes and Claims
2. Configuration Maps and Secrets
3. Stateful Sets






This Module (M5)

Table of Contents

1. Static Pods and Multi-container Pods
2. Autoscaling and Scheduling
3. Daemon Sets and Jobs
4. Ingress Resources and Controllers



Lab Infrastructure





Static Pods

Static Pods

- Static Pods are managed directly by the **kubelet**
- The API server is not looking after them
- The **spec** of a static Pod **cannot refer to other API objects**
- Their manifests are **standard** but are stored in a **specific folder**
- Usually, this is **/etc/kubernetes/manifests**
- And it is regulated by the **kublet configuration file**
- We can serve static Pods from the **local filesystem** or the **web**




Multi-Container Pods

Multi-container Pods

- Pods often have just **one container**
 - However, we may want do add **more than one**
 - This may be due to the **need of a helper process, or container with the same lifecycle**, etc.
 - There are three common **design patterns** for this
 - **Sidecar**
 - **Adapter**
 - **Ambassador**
- Single node patterns
(yes, there are multi node patterns as well)


Sidecar Pattern

- **Sidecar** containers enhance the main container
- For example, it may sync the local file system with a remote repository
- Or it may parse the logs of the main container and send them somewhere
- In any case, both share the same filesystem



Adapter Pattern

- **Adapter** containers are used to standardize and normalize the output
- This may be done in order to prepare it for a monitoring system
- This way, no matter the actual application or applications, the monitoring system will receive prepared data flow



Ambassador Pattern

- Ambassador containers proxy a local connection to the world
- The main container is connecting to a port on the localhost
- The ambassador container proxies the connection to the appropriate target
- Usually, this is used for providing access to a database
- For example, local one (when in test/dev) or a remote one (when in prod)





Init Containers

The Special Case of Init Containers

- **Specialized containers that run before app containers** in a Pod
- Contain utilities or setup scripts not present in an app image
- **App Containers** are specified via **containers** section and the **Init Containers** are specified via **initContainers** section
- Init Containers **always run to completion**. If one of them **fails**, the **kubelet** repeatedly **restarts it until it succeeds**
- Each Init Container **must complete successfully** before the **next one starts**

Native Sidecar Containers

- Introduced in Kubernetes 1.28
- In short, this is a **restartable init container**
- Available when the **SidecarContainers feature gate** is enabled (default since 1.29)

```
apiVersion: v1
kind: Pod
spec:
  initContainers:
    - name: secret-fetch
      image: secret-fetch:1.0
    - name: network-proxy
      image: network-proxy:1.0
  restartPolicy: Always
  containers:
    ...
```



Practice

Live Exercise in Class (Lab)



Autoscaling

Autoscaling

- Environments are not static but dynamic and changing
- This applies to the running pods and the resources they need
- Kubernetes, being a container orchestrator, has an answer. It offers the capability to perform autoscaling of resources
- There are three types
 - **Scale out the pods** by increasing their replica count
 - **Scale up the pods** by increasing their resources limits
 - **Scale out the cluster** by increasing the number of nodes

Scale Out Pods (Horizontal Pod Autoscaler)

- The **Horizontal Pod Autoscaler (HPA)** automatically scales the number of pods in a replication controller, deployment, replica set or stateful set based on observed CPU utilization
- It is implemented as a Kubernetes API **resource** and a **controller**
- The resource determines the behavior of the controller
- The controller periodically adjusts the number of replicas in a replication controller or deployment
- The Horizontal Pod Autoscaler is implemented as a **control loop**, with a period controlled by a flag with default value set to **15 seconds**
- Both upscale and downscale intervals are also controlled by flags and their default value is set to **5 minutes**

Scale Up Pods (Vertical Pod Autoscaler)



- The **Vertical Pod Autoscaler (VPA)** maintains the resource limits and requests for the containers in their pods up to date
- It can adjust the requests based on the usage. It also maintains the ratio between requests and limits
- Implemented via a **Custom Resource Definition (CRD)** object and has three components
 - **Recommender** monitors current and past resource consumption and provides recommended values
 - **Updater** checks which resources have correct resources set and if not, kills them in order to be recreated with updated values
 - **Admission Plugin** sets the correct resource requests on new pods

Scale Out Cluster Nodes (Cluster Autoscaler)



- Like the HPA but for cluster nodes
- Based on cluster utilization it can **change the number of nodes**
- This is useful also for **cost optimization**
- It checks for pods that cannot be scheduled on existing nodes. Then checks if node addition will solve the issue
- In the same manner, if pods can be rescheduled on other nodes to utilize them better, they will be evicted from a node and then the node will be removed



Scheduling

Taints and Tolerations

- **Taints** are applied to nodes and allow them to repel pods
- They have **key**, **value** and taint **effect** and are set like
kubectl taint nodes node1 key1=value1:NoSchedule
- Effect must be **NoSchedule** , **PreferNoSchedule** or **NoExecute**
- **Tolerations** are applied to pods and allow them to schedule on nodes with matching taints
- They are specified with **key**, **operator** (**Exists** or **Equal**), **value** (if the operator is equal) and **effect**

Node Selectors and Others

- A pod can be **constrained** to run only on **particular nodes**
- One of the ways to do this is to use **node selectors**
 - **nodeSelector** is a field part of the pod specification
 - It specifies a **map of key-value pairs**
 - For a pod to be able to run on a node, the node must have **all indicated key-value pairs** (it may have others as well)
- Alternatively, we can use **nodeName**, which select exact node
- Or we can use **(anti-)affinity** and **pod topology spread constraints**



Daemon Sets and Jobs

Workload Resources

- So far, we covered the following high level workload resources
 - **ReplicationController**
 - **ReplicaSet**
 - **Deployment**
 - **StatefulSet**
- There are a few others that we will cover now
 - **DaemonSet**
 - **Job and CronJob**

Daemon Sets

- Ensures that all (or some) Nodes run a copy of a Pod
- As nodes are added to the cluster, Pods are added to them
- As nodes are removed from the cluster, those Pods are garbage collected
- Deleting a **DaemonSet** will clean up the Pods it created

Jobs

- A **Job** creates one or more Pods and ensures that a specified number of them successfully terminate
- As pods successfully complete, the Job tracks the successful completions
- When a specified number of successful completions is reached, the Job is complete
- Deleting a Job will clean up the Pods it created
- The Job object will start a new Pod if the first Pod fails or is deleted
- They can run Pods either in **sequence** or in **parallel**

- A **CronJob** creates Jobs on a repeating schedule
- One CronJob object is like one line of a ***crontab*** (cron table) file
- CronJobs are useful for creating periodic and recurring tasks, like running backups, reports generation, sending emails, etc.



Practice

Live Exercise in Class (Lab)



Ingress Resources and Controllers


- **Ingress** exposes **HTTP** and **HTTPS** routes from outside the cluster to services within the cluster
- **Traffic routing** is controlled by **rules** defined on the Ingress resource
- We must have an **Ingress controller** to satisfy the Ingress
- Types
 - Single service (default backend)
 - Fanout
 - Name based virtual hosting
 - TLS

Ingress Controllers

- Ingress controllers are **not started automatically** with a cluster
- Kubernetes as a project supports and maintains **AWS, GCE, and nginx** ingress controllers
- Some of the others include **HAProxy, Istio, Contour**, etc.
- We may deploy **any number of ingress controllers** within a cluster
- Then, when we create one, we must annotate it with the appropriate **Ingress Class**


Fanout

- Routes traffic from a single IP address to more than one Service, based on the HTTP URI being requested



Name Based Virtual Hosting

- Supports routing HTTP traffic to multiple host names at the same IP address



Questions?



SoftUni



Software
University



SoftUni
AI



SoftUni
Digital



SoftUni
Creative



SoftUni
Foundation



Софтуни
БУДИТЕЛ



Finance
Academy

SoftUni Diamond Partners



**SUPER
HOSTING
.BG**



VIVACOM

Trainings @ Software University (SoftUni)



- Software University – High-Quality Education, Profession and Job for Software Developers
 - softuni.bg, about.softuni.bg
- Software University Foundation
 - softuni.foundation
- Software University @ Facebook
 - facebook.com/SoftwareUniversity



Software
University



- This course (slides, examples, demos, exercises, homework, documents, videos and other assets) is **copyrighted content**
- Unauthorized copy, reproduction or use is illegal
- © SoftUni – <https://about.softuni.bg/>
- © Software University – <https://softuni.bg>

