

Observability and Troubleshooting

Health and Status Checks

Logging, Auditing, and Troubleshooting

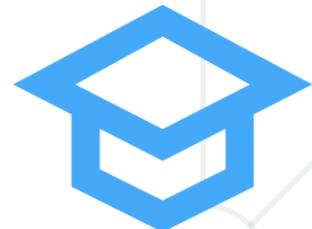


kubernetes

SoftUni Team

Technical Trainers

Software University

The Software University logo icon is an orange stylized 'S' shape.

SoftUni



Software University

<https://softuni.bg>

Have a Question?

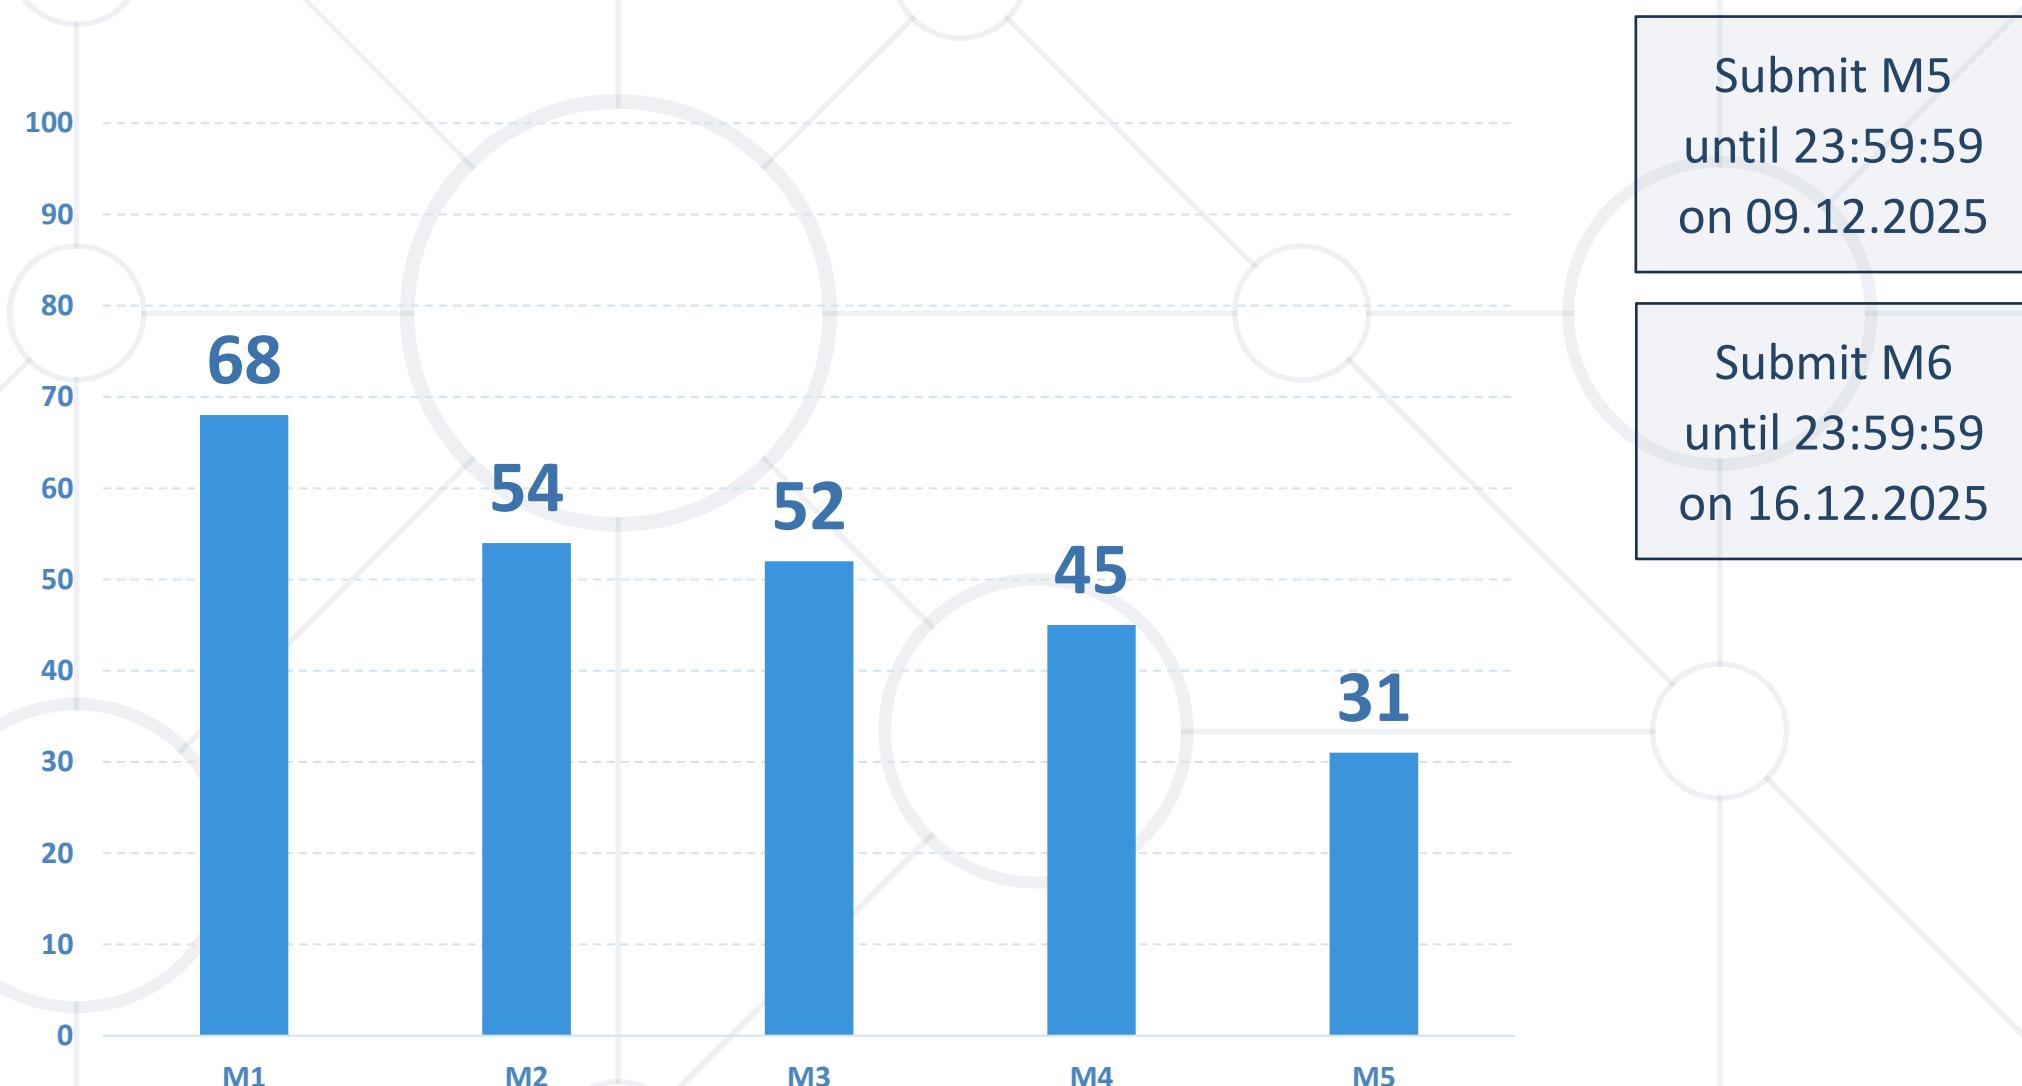


sli.do

#Kubernetes

[facebook.com
/groups/kubernetesnovember2025](https://facebook.com/groups/kubernetesnovember2025)

Homework Progress





Previous Module (M5)

Quick overview

Table of Contents

1. Static Pods and Multi-container Pods
2. Autoscaling and Scheduling
3. Daemon Sets and Jobs
4. Ingress Resources and Controllers





This Module (M6)

Table of Contents

1. Health and Status Checks
2. Auditing and Logging
3. Troubleshooting





Health and Status Checks

Health and Status Checks

- Periodic checks executed by the **kubelet** against containers
- Those checks are known as **probes**
- They can be **liveness**, **readiness**, and **startup** probes
- Their status can be either **Success**, **Failure**, or **Unknown**
- Used for a better control over the container and pod lifecycle and better integration with other objects

Check Methods

- Each probe type can use either **Exec**, **HTTP**, or **TCP** method
- **Exec** is used to **exec a specified command** inside the container.
It is considered **successful** if the **return code** is **0**
- **HTTP** makes a **GET request** against the pod's **IP address** on a **specified port** and **path**. It is considered **successful** if the **status code** is **between 200 and 399**
- **TCP** performs a check against the pod's **IP address** on a **specified port**. It is considered **successful** if the **port is open**

Liveness Probes (`livenessProbe`)

- Indicate whether a **container is running**
- If it fails, then **kubelet** kills the container
- After that, the container is **subject to the restart policy**
- It can be **Always**, **OnFailure**, and **Never**. The **default is Always**
- The **restart policy** is defined on **pod level** and applicable to all **containers** in the pod
- If no liveness probe is provided it is considered as if it was there and the return status is **Success**

Readiness Probes (readinessProbe)

- Indicate whether a container is **ready to respond to requests**
- If it **fails**, then the **endpoints controller removes** the pod's IP address from the **endpoints of all services** that match the pod
- A pod is considered ready when all its containers are ready
- The default state, before the initial delay is **Failure**
- If no readiness probe is provided it is considered as if it was there and the return status is **Success**

Startup Probes (startupProbe)

- Indicate whether the application in the container is **started**
- If it fails, then **kubelet** kills the container
- After that, the container is **subject to the restart policy**
- All **other probes** are **disabled** if a startup probe is present **until it succeeds**
- If no startup probe is provided it is considered as if it was there and the return status is **Success**

Common Fields

- **initialDelaySeconds** sets the number of seconds to wait before a probe to be initiated. **Defaults to 0 with minimal value of 0**
- **periodSeconds** sets how often (in seconds) a probe to be performed. **Defaults to 10 with minimal value of 1**
- **timeoutSeconds** sets the number of seconds before a probe times out. **Defaults to 1 with minimal value of 1**
- **successThreshold** sets the minimum consecutive successes for a probe to be considered successful after a failure. **Defaults to 1 with minimal value of 1**
- **failureThreshold** sets the number of times for Kubernetes to try failing probe before giving up (for **liveness – restart**, for **readiness – unready**). **Defaults to 3 with minimal value of 1**



Practice

Live Exercise in Class (Lab)



Auditing

Auditing

- Actions in the cluster are captured in chronological order
- They can be initiated by the users, applications, or control plane
- Answers who did what and when on what and what happened
- Audit records begin their existence in the **kube-apiserver**
- Each request on each stage of its execution generates an event
- It is pre-processed according to the policy and send to a backend
- Following stages are available **RequestReceived**, **ResponseStarted**, **ResponseComplete**, and **Panic**
- Audit logging may increase the memory consumption

<https://kubernetes.io/docs/tasks/debug-application-cluster/audit/>

Audit Policy

- Defines the rules about what events should be captured and what data they should include
- During processing, an event is compared against the list of rules in order
- First match sets the audit level of the event
- The available audit levels are **None**, **Metadata**, **Request**, and **RequestResponse**
- A policy to be valid, should have at least one rule

Audit Backends

- Audit backends persist audit events to an external storage
- Two backends are supported by default
- **Log backend** writes events into the filesystem
 - Writes audit events to a file in **JSONlines** format
 - Requires two volumes and volume mounts
- **Webhook backend** sends events to an **external HTTP API**
 - Both require **kube-apiserver** flags to be configured



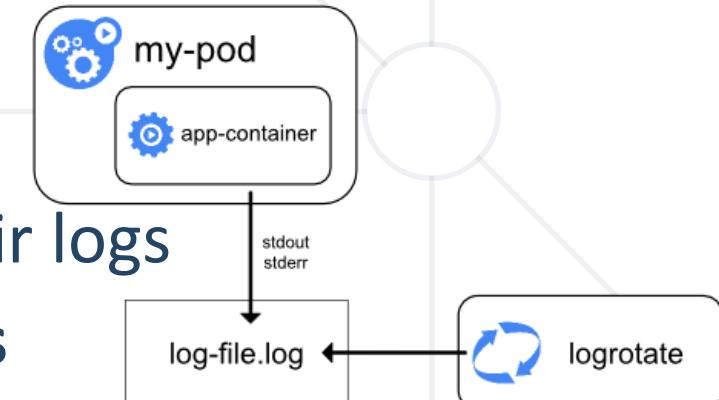
Logging

Logging Architecture

- Logs help us understand what is happening in our applications and cluster
- They are used for debugging problems and monitoring activity
- Most applications use logging either on the **stdout/stderr** or in a file
- Container engines/runtimes even though providing logging capabilities are usually not enough
- We need to access the logs even if and after a container or node crashes
- Thus, we need a **cluster-level logging** solution that will store logs elsewhere and they will have different lifecycle compared to the resources or nodes in the cluster

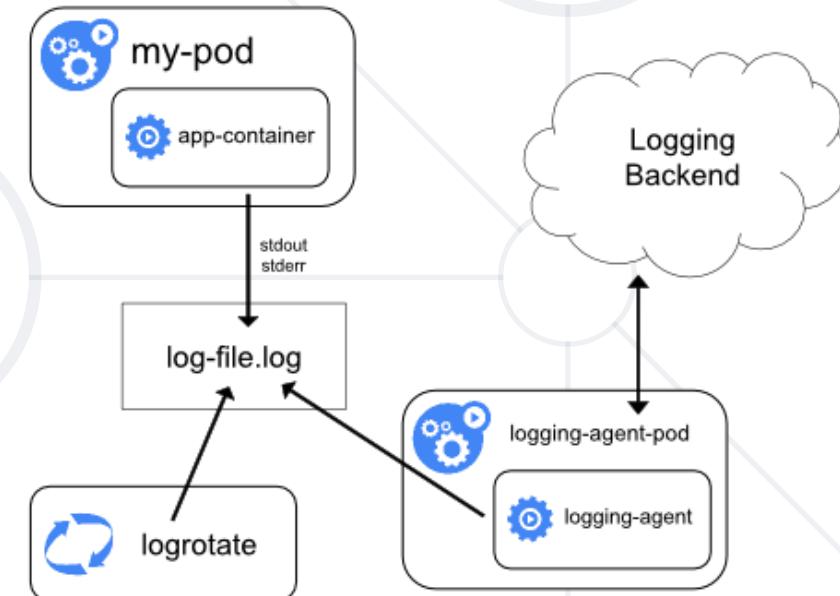
Node Level

- Even though not an ideal solution it can do the job
- We should pay attention to the following:
 - When a container is restarted, the **kubelet** keeps **one terminated container with its logs**
 - If a pod is evicted, all corresponding containers and their logs are deleted
 - We should set log rotation to do some housekeeping
 - Different container runtimes may have different requirements and capabilities
 - Not all system components are the same, so do their logs
 - Service based components log via **systemd routines**
 - Container based components use files in **/var/log**



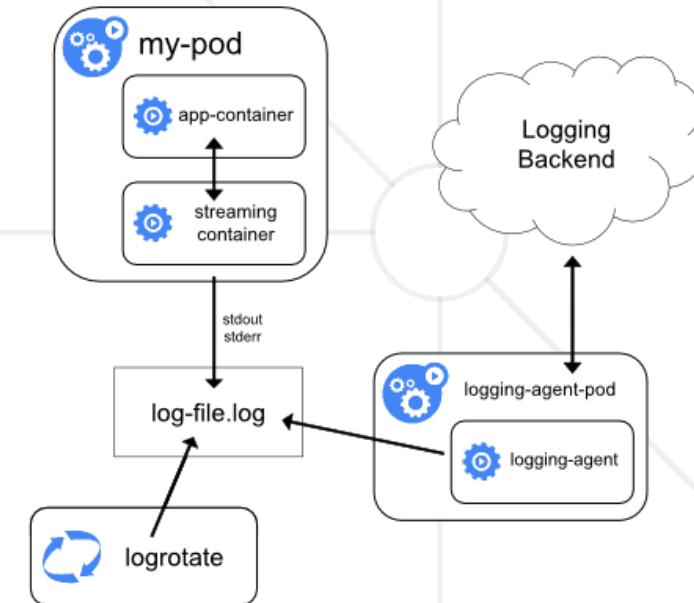
Node Logging Agent

- This is considered **cluster-level** logging approach
- For this, we deploy a node logging agent on each node
- Typically, the logging agent is containerized and deployed via **DaemonSet**
- The agent exposes the logs and pushes them to a backend



Streaming Sidecar and Logging Agent

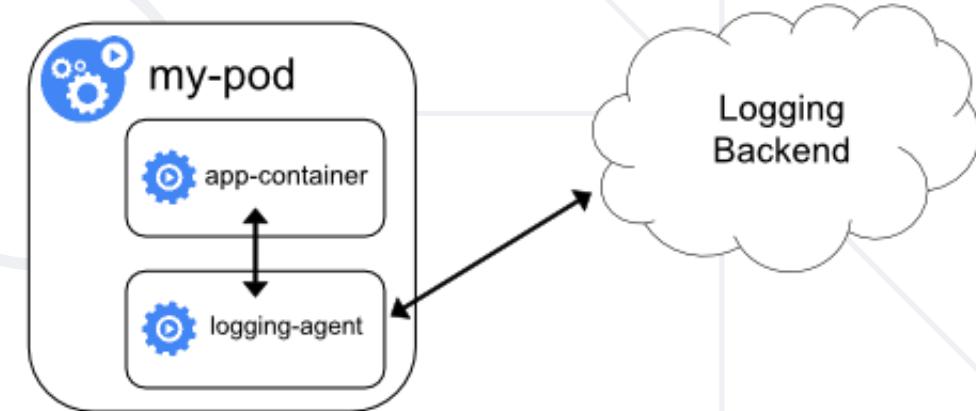
- This is considered **cluster-level** logging approach
- The sidecar container is publishing the log to its **stdout/stderr** and thus making it available for the **kubectl log** command
- Used to overcome limitations like separating multiple logs
- We can have more than one sidecar container



<https://kubernetes.io/docs/concepts/cluster-administration/logging/#streaming-sidecar-container>

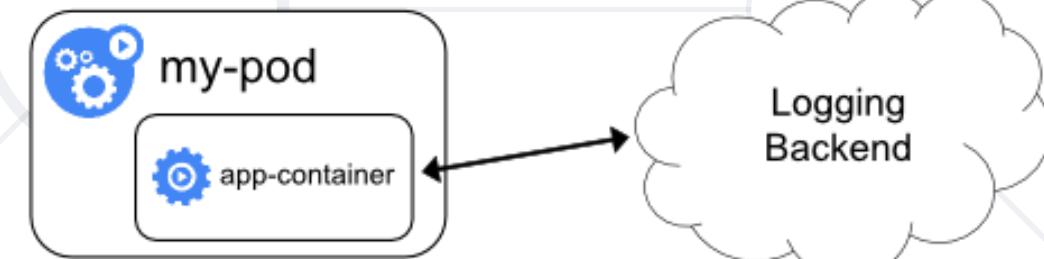
Sidecar with Logging Agent

- This is considered **cluster-level** logging approach
- Used when the node-level logging agent doesn't agree well with the application
- For this, we create a sidecar container with a logging agent that is especially configured and adjusted to the application's needs
- These logs are not consumable by the **kubectl log** command



Exposed Directly from the Application

- This is considered **cluster-level** logging approach
- Every application pushes its logs to a backend
- Simple solution which requires every application to support the common backend which may not always be feasible





Practice

Live Exercise in Class (Lab)



Troubleshooting

Troubleshooting

- Not always everything is going according to plan and things break
- The process of troubleshooting in Kubernetes is not that much different from other complex platforms
- Here, we have two distinct domains – **cluster** and **application**
- We should always check first the release notes for our version
- Once we narrowed down the cause, we should start applying the corrective measures one at a time until the issue is resolved

Cluster Troubleshooting *

- Common reasons include node power state, network connectivity, software version misalignment, data loss, bad configuration, etc.
- First, we should check if all nodes are there and operational
- Then, we must check the logs of the system components on the control plane and the workers
- Keep in mind that depending on how a component is deployed (native service or container) the logs may be in different places

Application Troubleshooting *

- First, we should define where exactly is the problem
- Is it the service
 - It is not reachable
 - It is not returning what is expected
- Or is it the workload object
 - What level and type of object
 - What is its state

Troubleshooting Pods

- Start with describing the pod
- Check the state of all containers inside the pod
- If it is in **pending** state, then it cannot be scheduled on a node.
Usually, this is because of **lack of resources**
- If it is in **waiting** state, then it is scheduled, but still cannot run.
Usually, this is because of a **wrong or missing image**
- It is in **running** state but doesn't behave as expected. Usually,
this is because of a **manifest error**

Troubleshooting Services *

- Does the service exist
- Is it defined correctly
- Are there any endpoints at all and how many
- Are those pods working
- Is the service reachable by DNS name and/or IP
- Is the kube-proxy working



Practice

Live Exercise in Class (Lab)

Questions?



SoftUni

Software
UniversitySoftUni
AISoftUni
DigitalSoftUni
CreativeSoftUni
FoundationСофтуни
БУДИТЕЛFinance
Academy

SoftUni Diamond Partners



**SUPER
HOSTING
.BG**



VIVACOM

Trainings @ Software University (SoftUni)



- Software University – High-Quality Education, Profession and Job for Software Developers
 - softuni.bg, about.softuni.bg
- Software University Foundation
 - softuni.foundation
- Software University @ Facebook
 - facebook.com/SoftwareUniversity



Software
University



- This course (slides, examples, demos, exercises, homework, documents, videos and other assets) is **copyrighted content**
- Unauthorized copy, reproduction or use is illegal
- © SoftUni – <https://about.softuni.bg/>
- © Software University – <https://softuni.bg>

