

Security and Policies

Authentication, Authorization, and Admission.

Resource Management. Network Policies



kubernetes

SoftUni Team

Technical Trainers

Software University

The logo icon for Software University is an orange stylized 'S' shape.

SoftUni



Software University

<https://softuni.bg>

Have a Question?

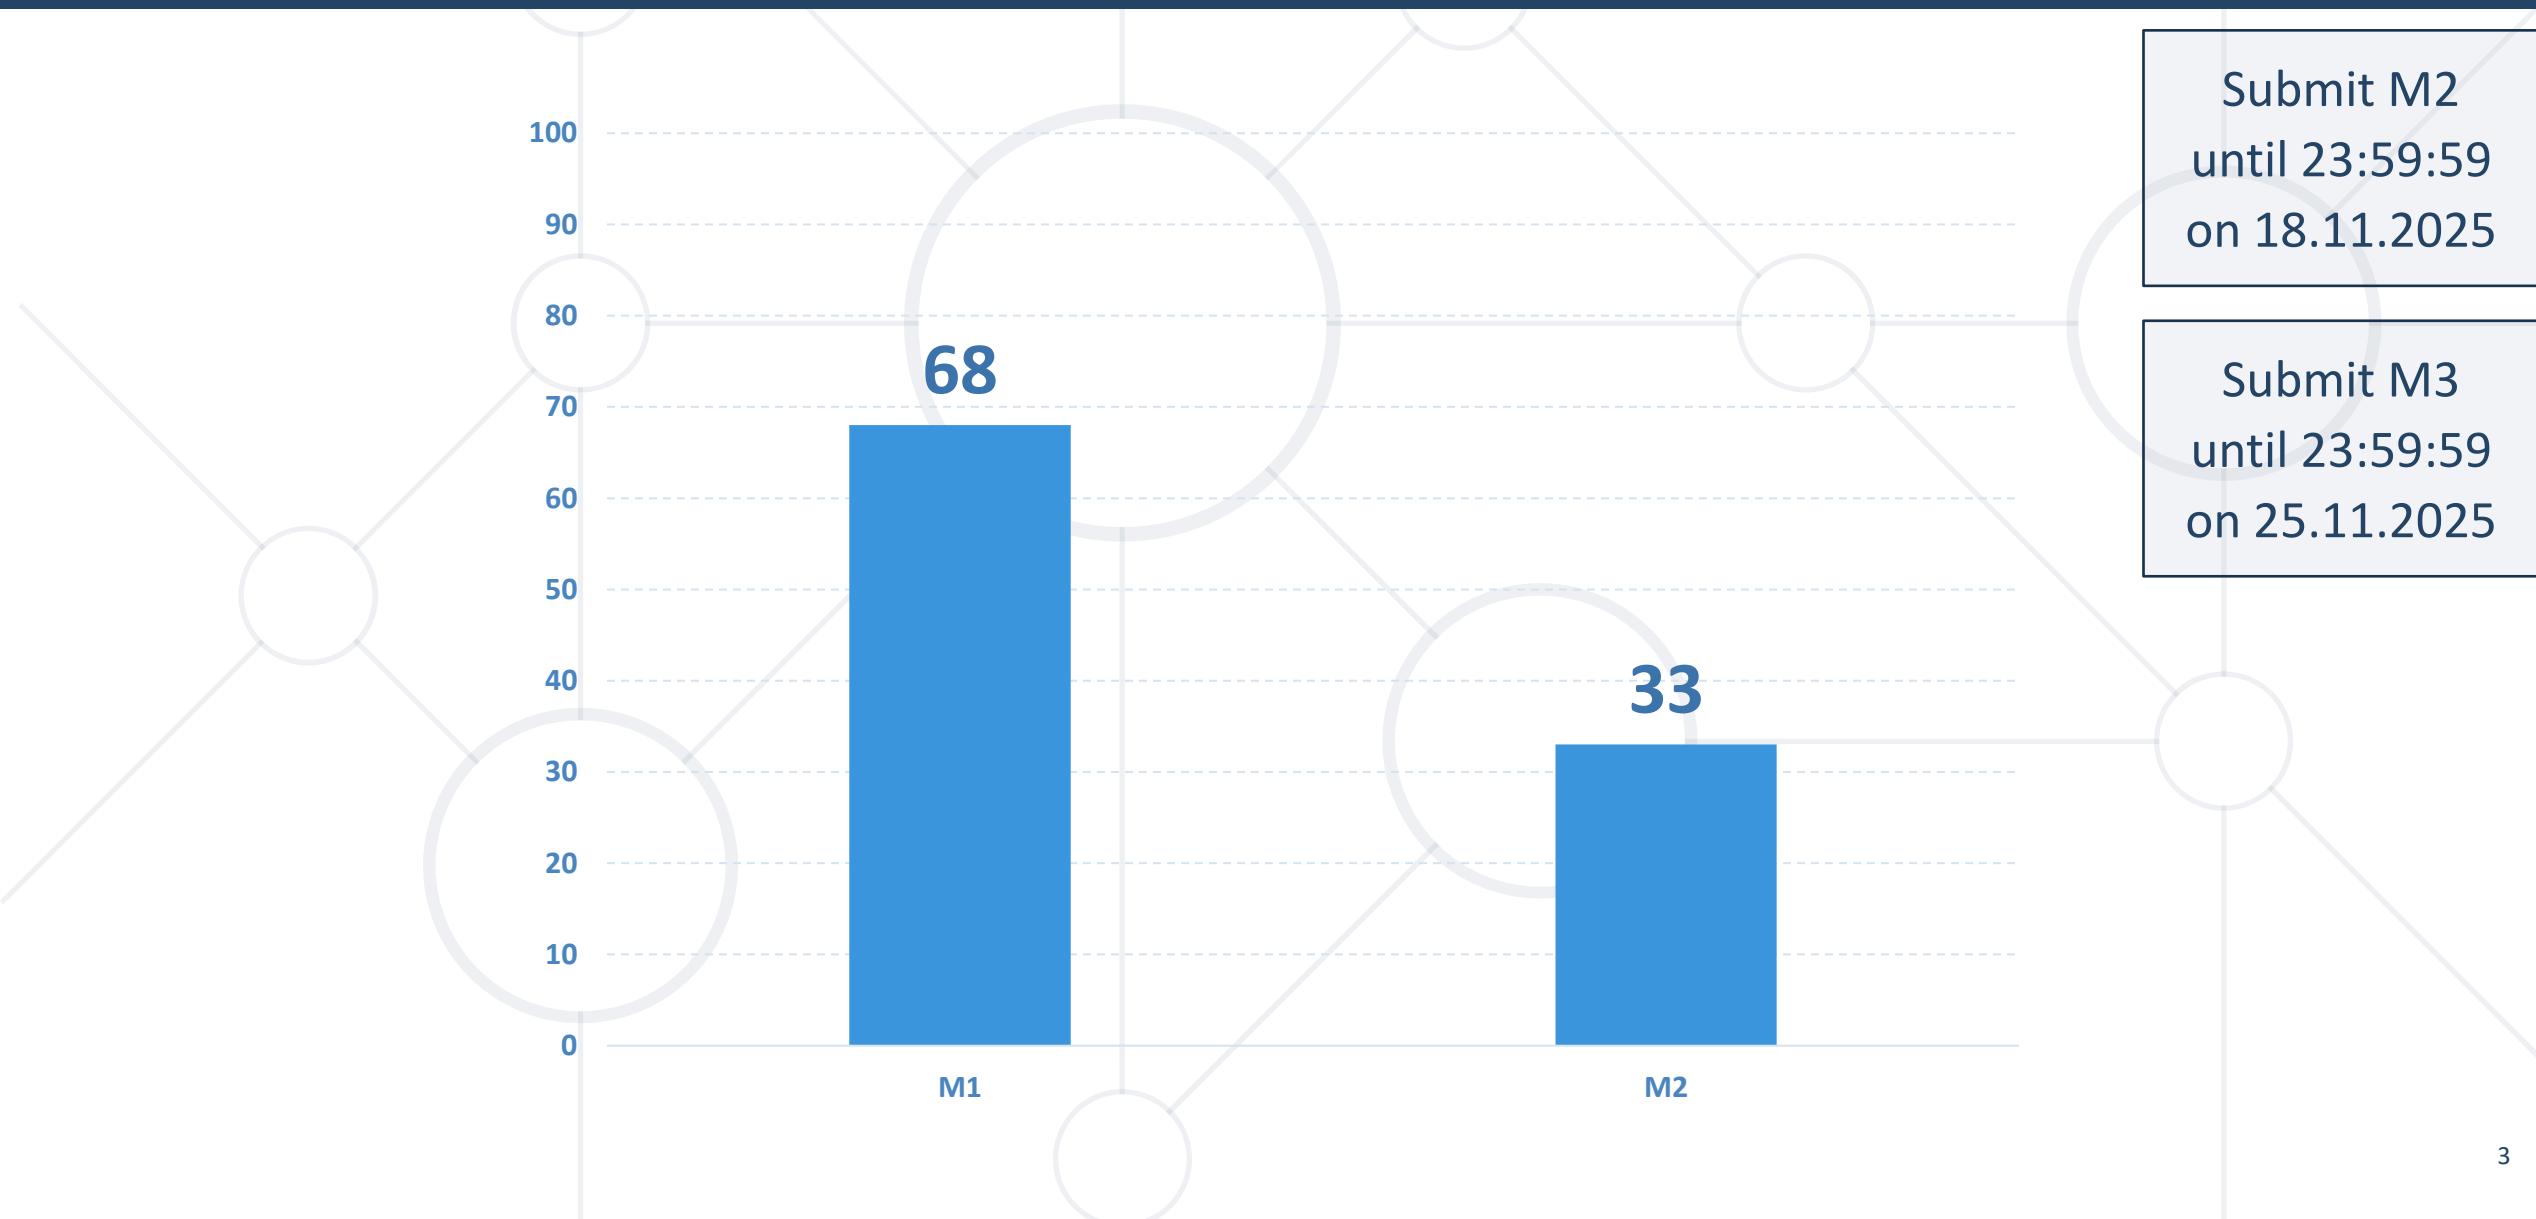


sli.do

#Kubernetes

[facebook.com
/groups/kubernetesnovember2025](https://facebook.com/groups/kubernetesnovember2025)

Homework Progress (as of 17:00)





Previous Module (M2)

Quick overview

Table of Contents

1. Basic Cluster Installation
2. Manage and Upgrade Kubernetes Cluster
3. Highly-available Kubernetes Cluster





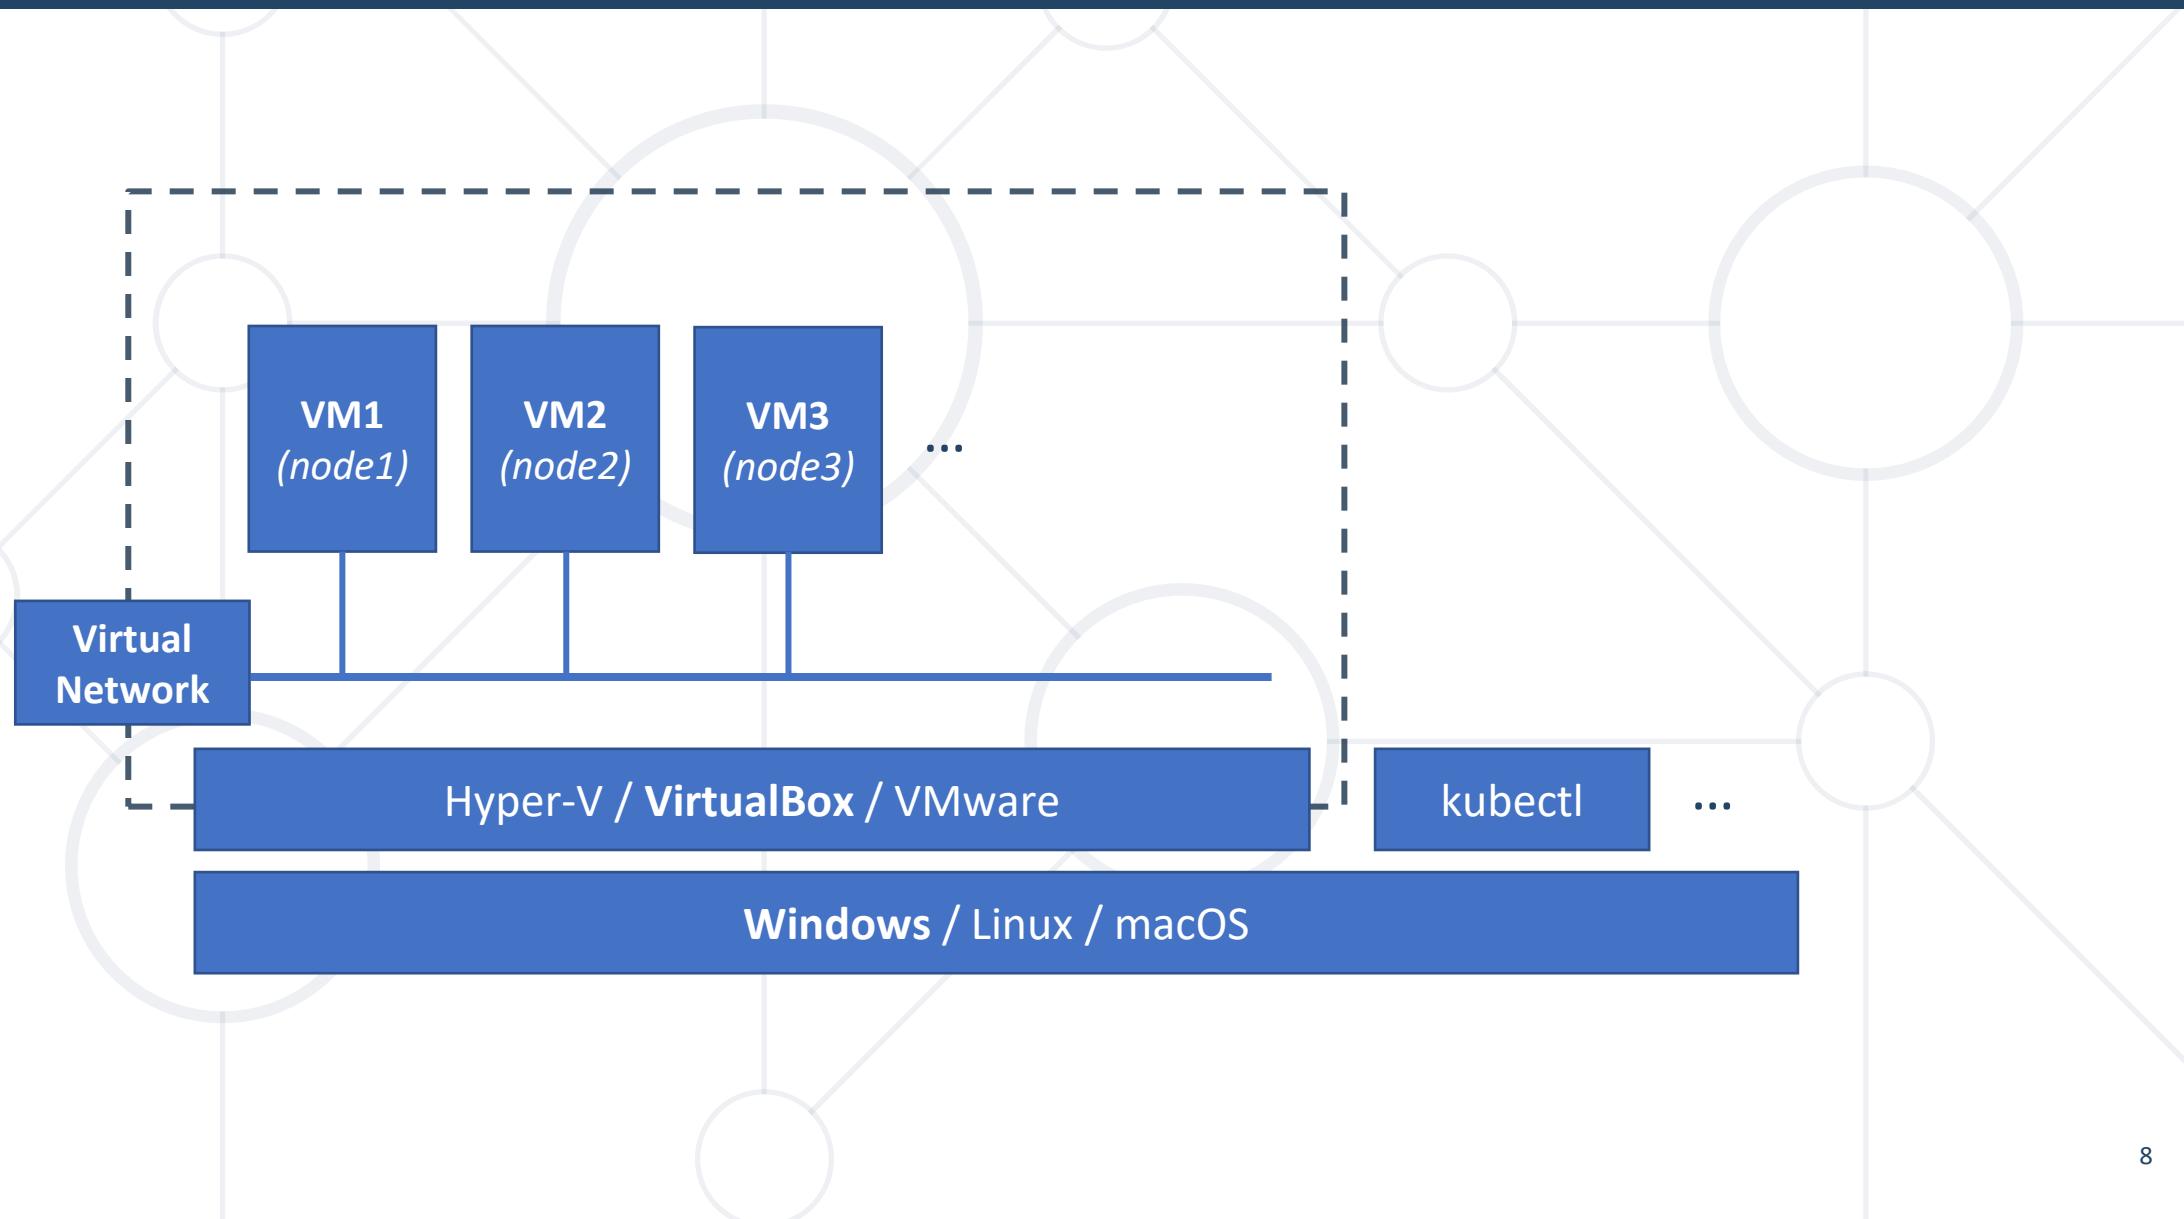
This Module (M3)

Table of Contents

1. Could Native Security
2. Authentication, Authorization and Admission Control
3. Resource Requirements, Limits and Quotas
4. Network Policies



Lab Infrastructure

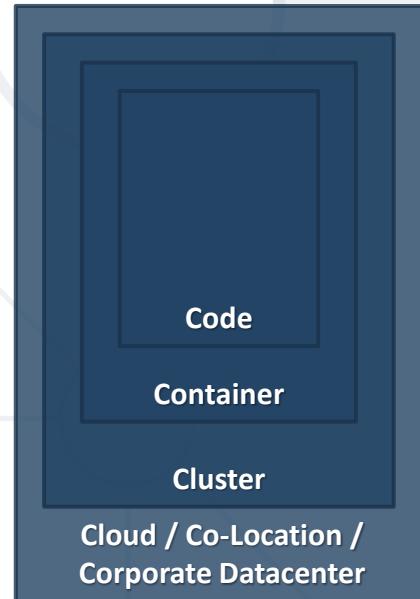




Cloud Native Security

The 4C's of Cloud Native *Kubernetes* Security

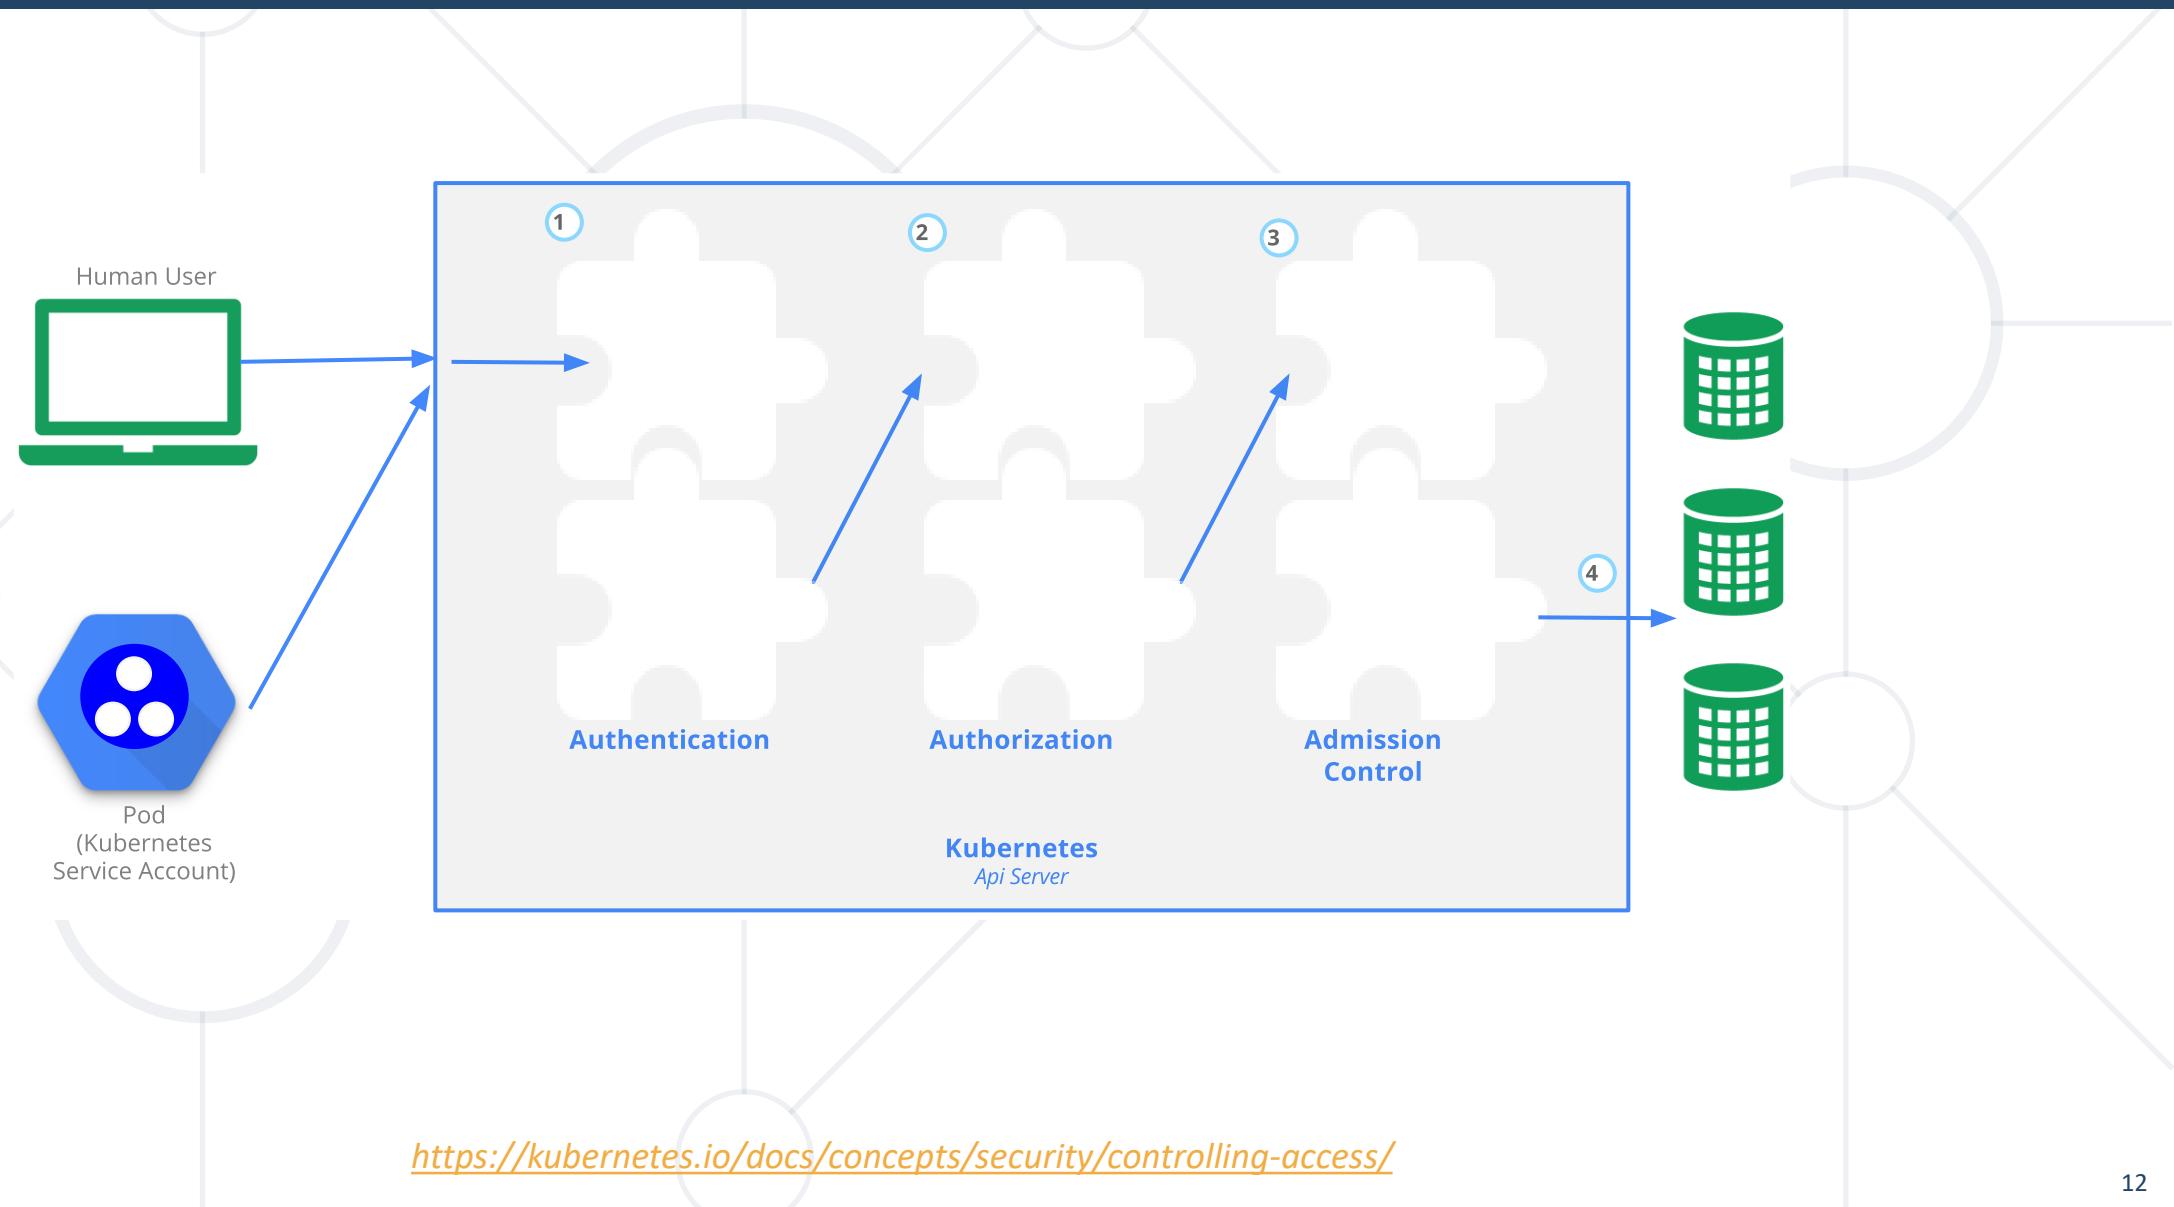
- **Cloud / Co-Location / Corporate Datacenter**
 - General: Least privilege, network access, encryption, misconfiguration, default settings, etc.
 - Kubernetes: Access to nodes and API server, access and encryption of etcd
- **Cluster**
 - General: Cluster components and applications
 - Kubernetes: RBAC, Authentication, Secrets, Network Policies, etc.
- **Container**
 - Secure base images, image signing, scan for vulnerabilities, container runtime
- **Code**
 - Static code analysis, 3rd party dependencies, access over TLS, exposed ports, hardcoded secrets, etc.





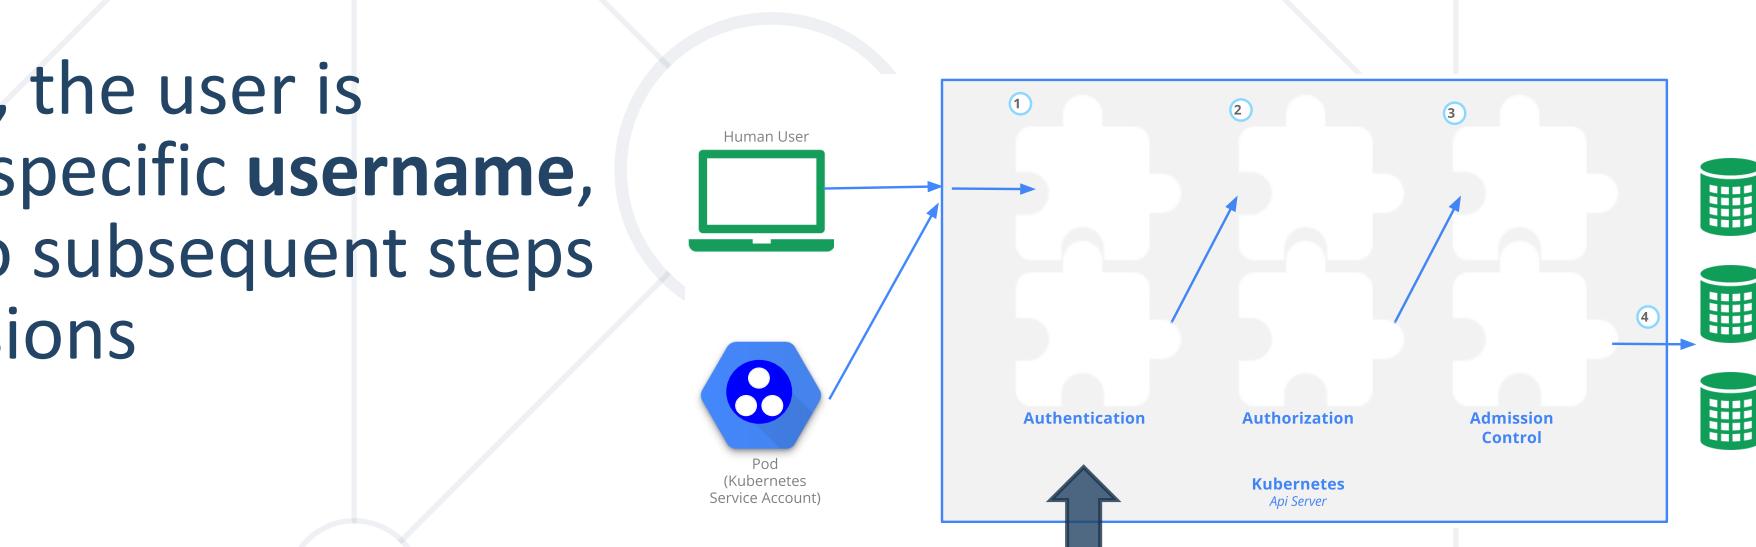
Authentication, Authorization and Admission Control

Access Control Overview



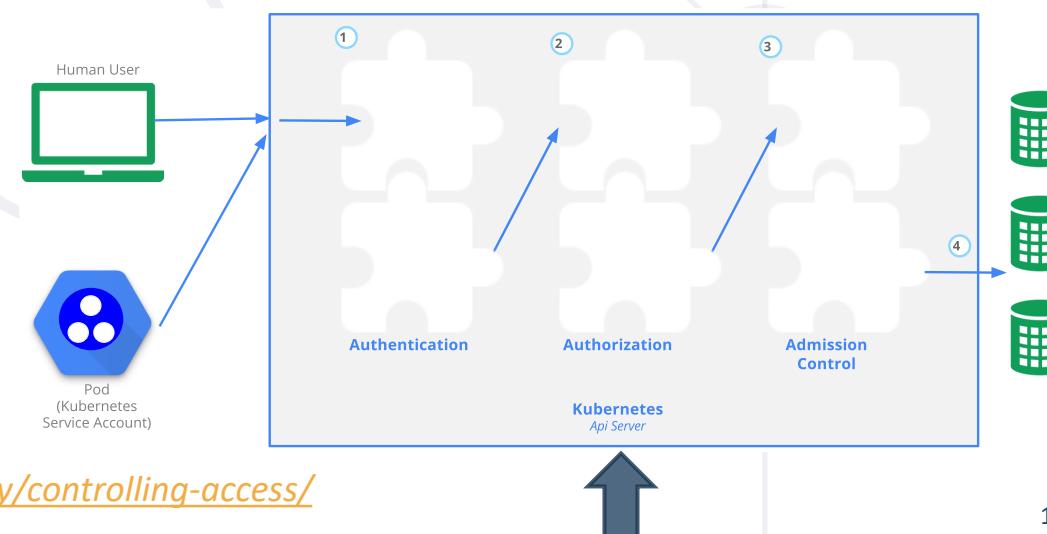
Authentication

- Authentication modules include **client certificates, password, and plain tokens, bootstrap tokens, and JSON Web Tokens**
- Multiple authentication modules can be specified. Each one is tried in sequence, until one of them succeeds
- If the request cannot be authenticated, it is rejected with HTTP status **code 401**
- If process succeeds, the user is authenticated as a specific **username**, and it is available to subsequent steps to use in their decisions



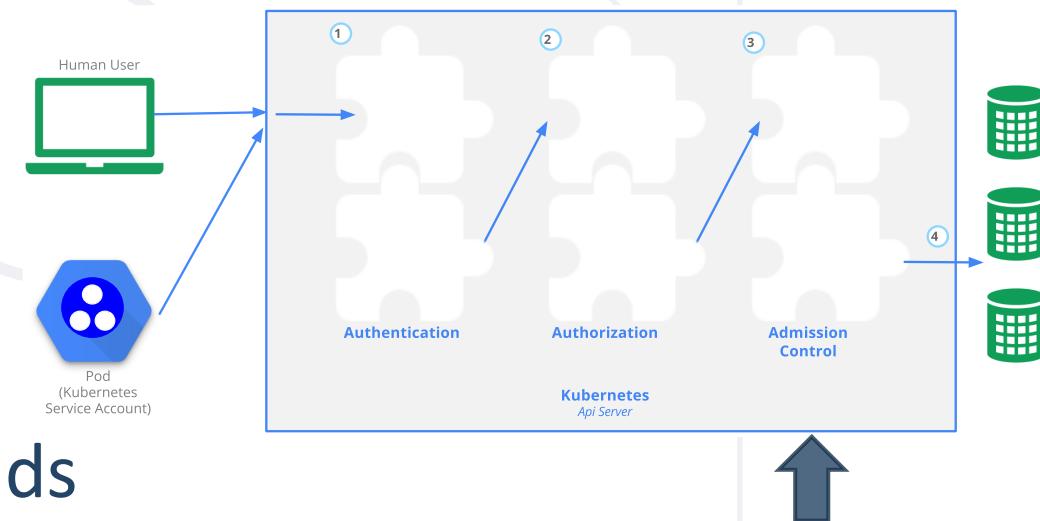
Authorization

- After a successful authentication, the request must be authorized
- It must include the **username**, the **action**, and the **target object**
- It is authorized if an existing policy confirms that the user has the right to complete the requested action
- Multiple authorization modules are supported, such as **ABAC**, **RBAC**, and **Webhook**
- They are configured during the cluster creation
- More than one may be specified
- All are evaluated until one authorizes the request
- If all deny it, then the request is denied (**403**)



Admission Control

- Software modules that can **modify or reject** requests
- In addition to the attributes available to Authorization modules, Admission Control modules can access the **contents of the object** that is being created or modified
- **Multiple controllers** may be configured. In this case they are **called in order**
- If any admission controller module rejects, then the request is immediately rejected
- They act on requests that **create, modify, delete, or connect to (proxy)** an object
- They **ignore** requests that merely **read objects**
- They can also set complex defaults for fields



User Accounts vs Service Accounts

- When we (humans) access the cluster (for example with kubectl), we are authenticated by the **apiserver** as a particular **User Account** (usually **admin**)
- Processes in containers inside pods can also contact the **apiserver**. When they do, they are authenticated as a particular **Service Account** (for example, **default**)

Service Accounts

- Each namespace gets one by default named **default**
- During pod creation we can specify a service account
- If we omit it, the default one is assigned
- Service account names are formatted like this:
system:serviceaccount:<namespace>:<service account name>
- We can create additional ones to further adjust pods' rights
- Pods can use service accounts from the same namespace
- Each pod is attached to only one service account
- One service account may be shared by multiple pods

Role-based Access Control (RBAC)

- Controls who can do what with which type of resources
- **Roles** contain rules that represent a set of permissions
- Permissions are purely additive (there are no "deny" rules)
- A **role binding** grants the permissions defined in a role to a user or set of users
- Role bindings hold a list of **subjects** (users, groups, or service accounts), and a **reference to the role** being granted

Roles and Bindings

- **Roles** set permissions within a namespace
- **ClusterRoles** are non-namespaced resources. They are used to
 - Grant permissions on namespaced resources within individual namespaces
 - Grant permissions on namespaced resources across all namespaces
 - Grant permissions on cluster-scoped resources

Roles and Bindings

- **RoleBinding** grants permissions within a specific namespace whereas a **ClusterRoleBinding** grants that access cluster-wide
- RoleBinding may reference any Role in the same namespace or a ClusterRole
- If we want to bind a ClusterRole to all namespaces, we must use ClusterRoleBinding



Practice

Live Exercise in Class (Lab)



Resource Requirements, Limits and Quotas

Resource Requests and Limits

- By default, containers run with unbounded resources on a cluster
- We can control how much of a resource is granted to a Container
- Most common resources are the CPU and memory
- We specify how much resources a Container needs to operate via the **request** option. This information is accumulated on a Pod level and a decision for scheduling of the Pod is made
- We restrict how much of a resource a Container can use via the **limit** option

- The CPU resource is measured in **CPU units**
- One CPU, in Kubernetes, is equivalent to **1 AWS vCPU or 1 GCP Core or 1 Azure vCore or 1 Hyperthread on a bare-metal Intel processor with Hyperthreading**
- We can specify fractional values, for example **0.5** which is the half of 1 CPU
- The above can be stated as **500m** (milicpu) units

- Memory requests and limits are measured in **bytes**
- We can specify it using **integer** values or **fixed-point** numbers
- We may use the standard suffixes **k, M, G, T, etc.** or the power-of-two ones **Ki, Mi, Gi, Ti, etc.**
- For example, to set **128 MiB** of memory, we must use **128Mi**

- Resource quotas are a tool for administrators to address the need of sharing cluster resources between teams of users
- They are defined by the **ResourceQuota** object
- Provide constraints that **limit aggregate resource consumption per namespace**
- Limit the **quantity of objects** that can be created in a namespace by type
- Limit the **total amount of compute resources** that may be consumed by resources

Limit Ranges

- A **LimitRange** is a policy to constrain resource allocations (to **Pods** or **Containers**) in a namespace
- Enforce **minimum** and **maximum** compute resources usage per Pod or Container in a namespace
- Enforce minimum and maximum storage request per PersistentVolumeClaim in a namespace
- Enforce a ratio between request and limit for a resource in a namespace
- Set **default request/limit** for compute resources in a namespace and automatically inject them to Containers at runtime



Practice

Live Exercise in Class (Lab)



Network Policies

Network Policies

- Control traffic flow at the IP address or port level for particular applications in the cluster
- Allow control if and how a pod is allowed to communicate with various network entities over the network
- The entities are identified by a combination of
 - Other pods that are allowed
 - Namespaces that are allowed
 - IP blocks

Prerequisites

- They are **implemented by the network plugin**
- **Flannel** does not support Network Policies
- Plugins like **Antrea**, **Calico**, and **Cilium** do support them
- Creating a **NetworkPolicy** resource without a plugin that supports it won't have any effect

Main Elements

- **podSelector** selects a pod or group of pods. If an empty one is used this matches all pods in the namespace
- **policyTypes** can include either **Ingress** or **Egress** or both. It specifies which traffic for the selected pods is being regulated
- **ingress** includes a list of allowed ingress rules. Each rule consists of **from** and **ports** sections
- **egress** includes a list of allowed egress rules. Each rule consists of **to** and **ports** sections

<https://kubernetes.io/docs/concepts/services-networking/network-policies/>

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
  ingress:
    - from:
        - ipBlock:
            cidr: 172.17.0.0/16
        except:
          - 172.17.1.0/24
    - namespaceSelector:
        matchLabels:
          project: myproject
    - podSelector:
        matchLabels:
          role: frontend
  ports:
    - protocol: TCP
      port: 6379
```

Default Policies

- By default, no policies exist in a namespace, so all ingress and egress traffic is allowed
- Should we want, we can create the following **default network policies**

Deny all ingress traffic

```
---  
apiVersion: networking.k8s.io/v1  
kind: NetworkPolicy  
metadata:  
  name: default-deny-ingress  
spec:  
  podSelector: {}  
  policyTypes:  
    - Ingress
```

Allow all ingress traffic

```
---  
apiVersion: networking.k8s.io/v1  
kind: NetworkPolicy  
metadata:  
  name: allow-all-ingress  
spec:  
  podSelector: {}  
  ingress:  
    - {}  
  policyTypes:  
    - Ingress
```

Deny all egress traffic

```
---  
apiVersion: networking.k8s.io/v1  
kind: NetworkPolicy  
metadata:  
  name: default-deny-egress  
spec:  
  podSelector: {}  
  policyTypes:  
    - Egress
```

Allow all egress traffic

```
---  
apiVersion: networking.k8s.io/v1  
kind: NetworkPolicy  
metadata:  
  name: allow-all-egress  
spec:  
  podSelector: {}  
  egress:  
    - {}  
  policyTypes:  
    - Egress
```



Practice

Live Exercise in Class (Lab)

Questions?



SoftUni



Software
University



SoftUni
AI



SoftUni
Digital



SoftUni
Creative



SoftUni
Foundation



Софтуни
БУДИТЕЛ



Finance
Academy

SoftUni Diamond Partners



Trainings @ Software University (SoftUni)



- Software University – High-Quality Education, Profession and Job for Software Developers
 - softuni.bg, about.softuni.bg
- Software University Foundation
 - softuni.foundation
- Software University @ Facebook
 - facebook.com/SoftwareUniversity



Software
University



- This course (slides, examples, demos, exercises, homework, documents, videos and other assets) is **copyrighted content**
- Unauthorized copy, reproduction or use is illegal
- © SoftUni – <https://about.softuni.bg/>
- © Software University – <https://softuni.bg>

